

INTRODUCTION TO COMPUTER VISION

Lecture 6 – Depth Estimation

Gyeongsik Moon

Visual Computing and AI Lab

Korea University



**Visual
Computing
and AI Lab**

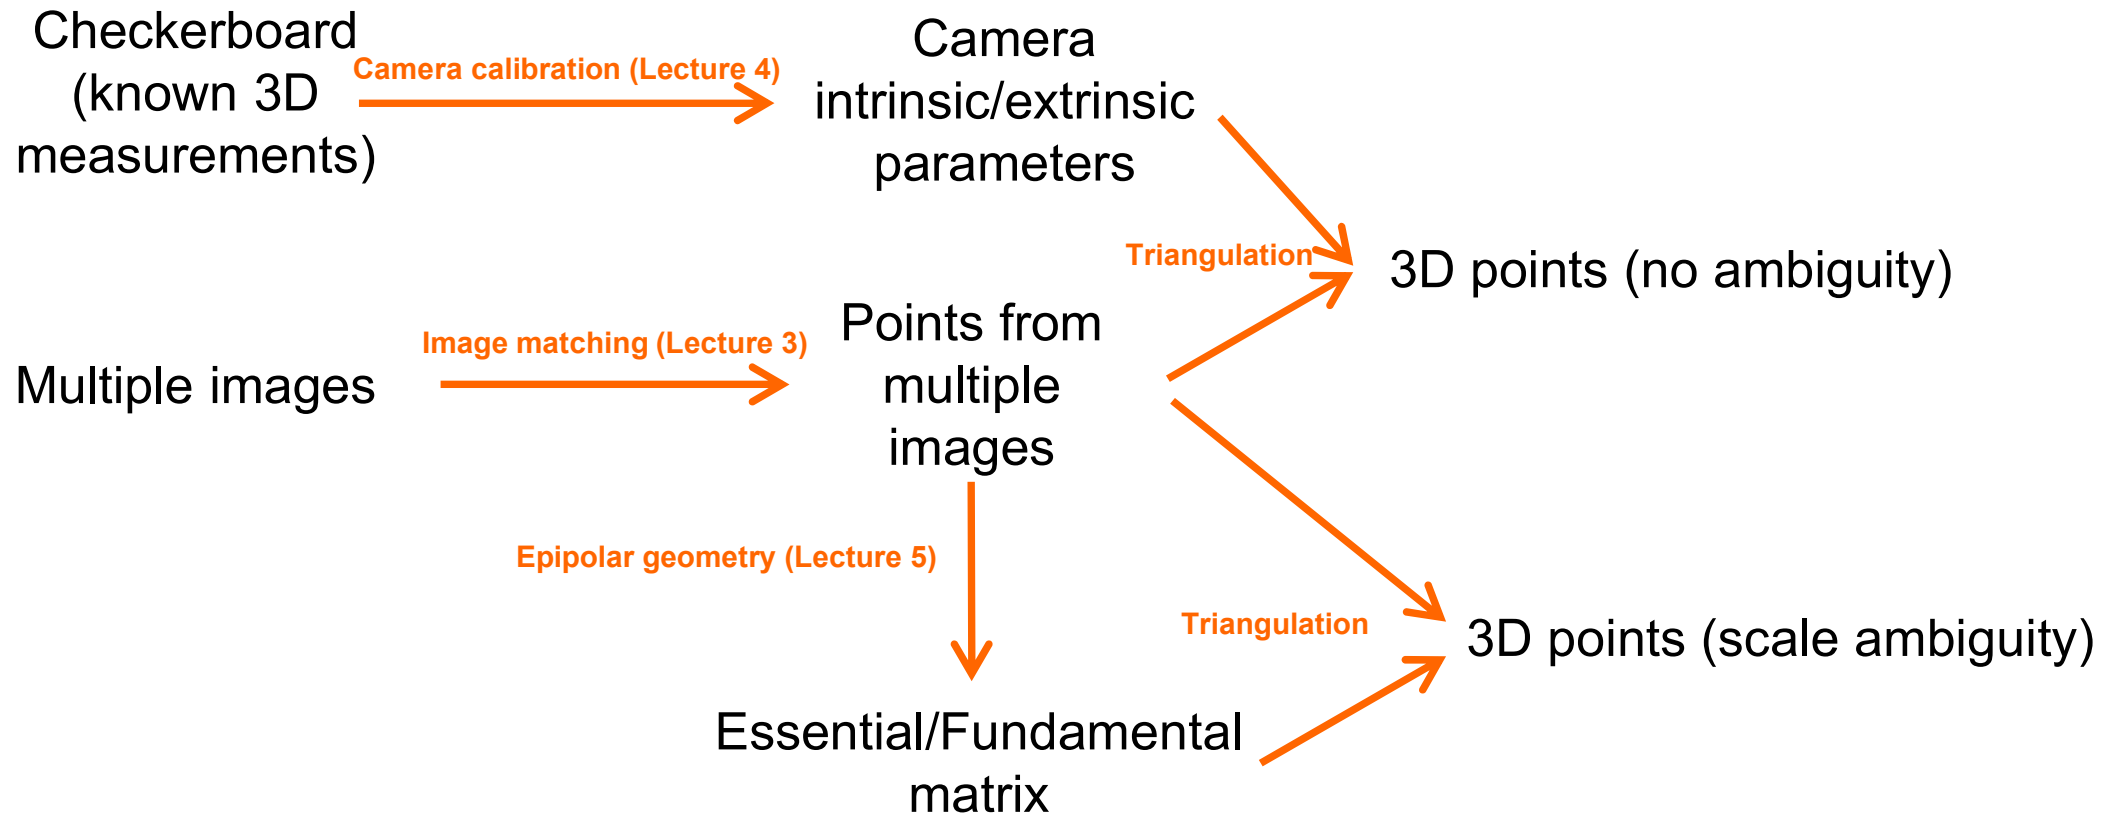
1. Triangulation

2. Bundle Adjustment

1. Triangulation

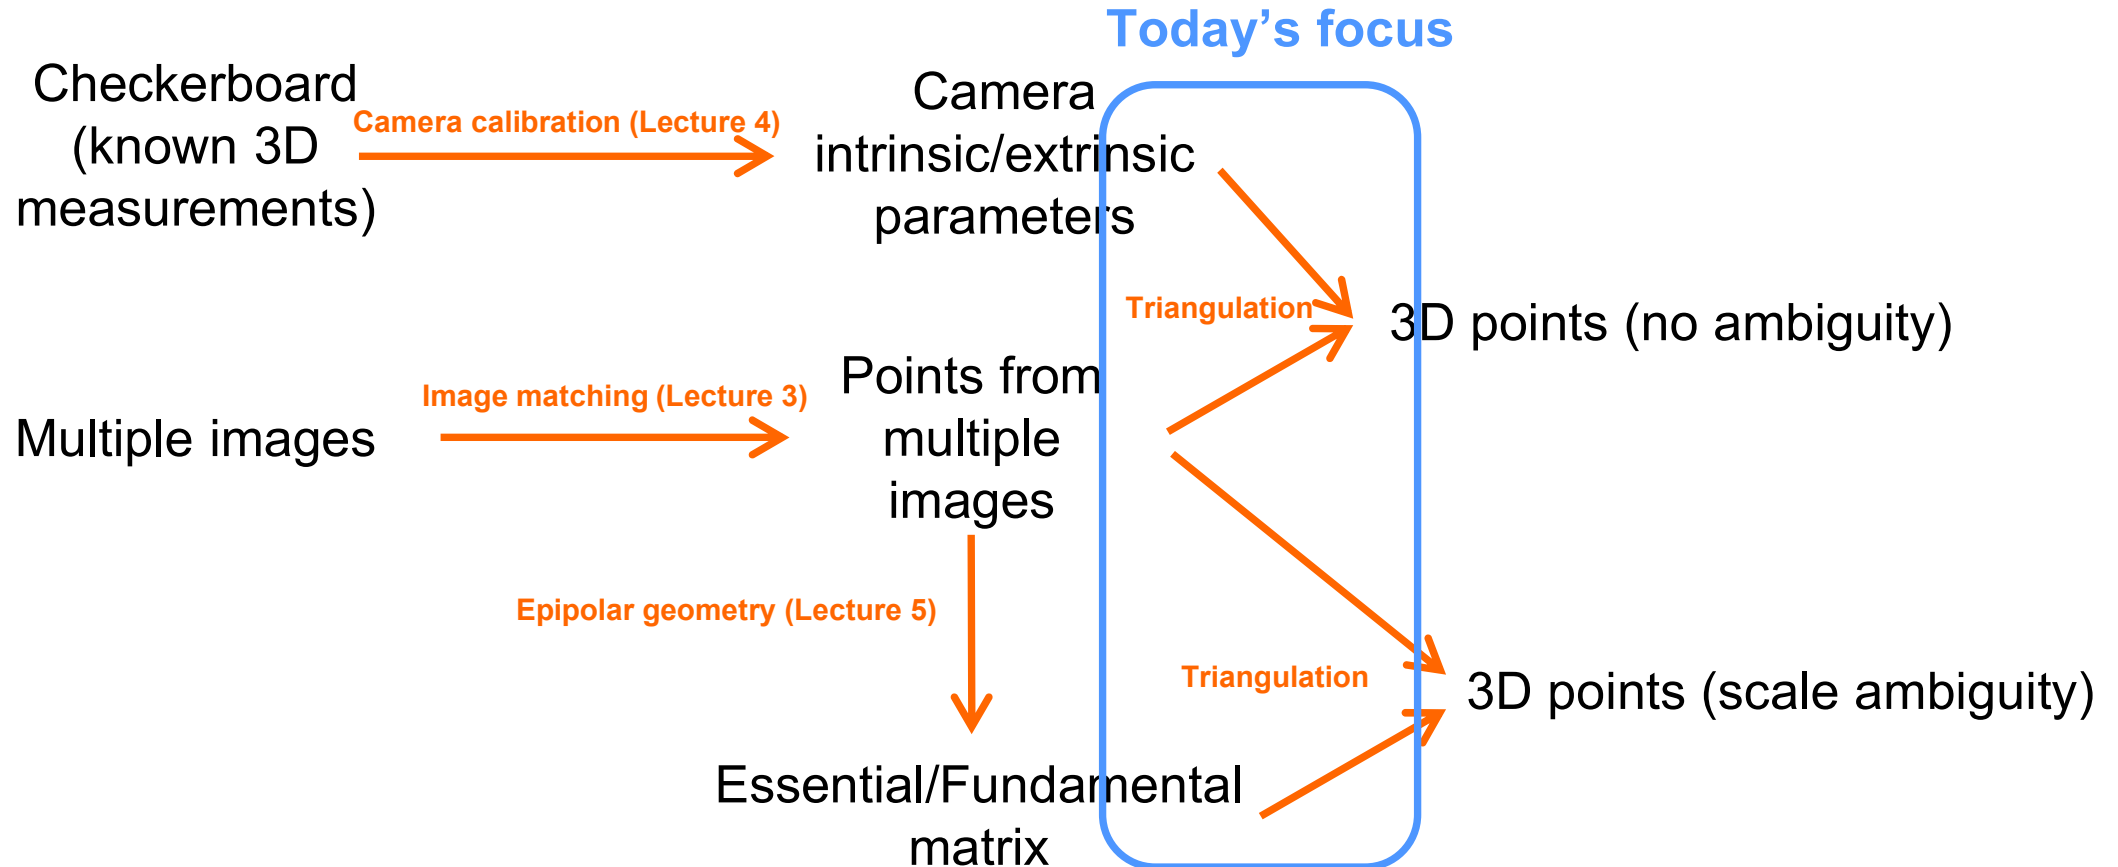
- All the entire pipeline: multi-view geometry theory

Slide from Lecture 5




- All the entire pipeline: multi-view geometry theory


Slide from Lecture 5




Triangulation



Triangulation



Triangulation



Triangulation

Given a set of (noisy) matched points

$$\{x_i, x'_i\}$$

2D matched point

and camera matrices

$$\mathbf{P}, \mathbf{P}'$$

intrinsic/extrinsic


Estimate the 3D point

$$\mathbf{X}$$

$$\overset{2D}{\mathbf{x}} = \mathbf{P} \overset{3D}{\mathbf{X}}$$

known known

Can we compute \mathbf{X} from a single correspondence \mathbf{x} ? *in theory, no unique point*



$$\mathbf{x} = \mathbf{P} \mathbf{X}$$

known

known

*Can we compute \mathbf{X} from two
correspondences \mathbf{x} and \mathbf{x}' ?*

$$\mathbf{x} = \mathbf{P} \mathbf{X}$$

known

known

*Can we compute \mathbf{X} from two
correspondences \mathbf{x} and \mathbf{x}' ?*

yes if perfect measurements

$$\mathbf{x} = \mathbf{P} \mathbf{X}$$

known

known

Can we compute \mathbf{X} from two correspondences \mathbf{x} and \mathbf{x}' ?

yes if perfect measurements

There will not be a point that satisfies both constraints
because the measurements are usually noisy

$$\mathbf{x}' = \mathbf{P}' \mathbf{X} \quad \mathbf{x} = \mathbf{P} \mathbf{X}$$

assumption: same 3D point.

Need to find the best fit

optimize

$$\mathbf{x} = \mathbf{P}X$$

(homogeneous
coordinate)

Also, this is a similarity relation because it involves homogeneous coordinates

$$\mathbf{x} = \alpha \mathbf{P}X$$

(inhomogeneous
coordinate)

Same ray direction but differs by a scale factor

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \alpha \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

How do we solve for unknowns in a similarity relation?

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

(homogeneous
coordinate)

Also, this is a similarity relation because it involves homogeneous coordinates

$$\mathbf{x} = \alpha \mathbf{P}\mathbf{X}$$

(inhomogeneous
coordinate)

Same ray direction but differs by a scale factor

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \alpha \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

How do we solve for unknowns in a similarity relation?


Direct Linear Transform (DLT)

Remove scale factor, convert to linear system and solve with SVD.

Comparison to DLT from Lecture 4

Tsai's Method: Calibration from 3D Objects Slide from Lecture 4

- This method was proposed in 1987 by Tsai and consists of measuring the 3D position of $n \geq 6$ **control points** on a 3D calibration target and the **2D coordinates of their projection** in the image.
- Assumption: we know 2D and 3D coordinates of control points
 - 2D: image pre-processing (e.g., corner detectors)
 - 3D: we know actual size of the 3D object and actual positions of control points as well



Tsai, Roger Y. (1987) "A Versatile Camera Calibration Technique for High-Accuracy 3D Machine Vision Metrology Using Off-the-Shelf TV Cameras and Lenses," *IEEE Journal of Robotics and Automation*, 1987. [PDF](#).

Applying the Direct Linear Transform (DLT) algorithm

The idea of the DLT is to rewrite the perspective projection equation as a **homogeneous linear equation** and solve it by standard methods. Let's write the perspective equation for a generic 3D-2D point correspondence:

$$\lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K[R | T] \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix} \Rightarrow$$

$$\Rightarrow \lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha_u & 0 & u_0 \\ 0 & \alpha_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

$$\Rightarrow \lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha_u r_{11} + u_0 r_{31} & \alpha_u r_{12} + u_0 r_{32} & \alpha_u r_{13} + u_0 r_{33} & \alpha_u t_1 + u_0 t_3 \\ \alpha_v r_{21} + v_0 r_{31} & \alpha_v r_{22} + v_0 r_{32} & \alpha_v r_{23} + v_0 r_{33} & \alpha_v t_2 + v_0 t_3 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

- ① camera parameter
- ② 2D matching
- ③ 3D point.

Comparison to DLT from Lecture 4

- Both are for solving linear equations
- Lecture 4: Getting camera parameters from 1) estimated 2D matching and 2) 3D measurements
- Today: Getting 3D points from 1) estimated 2D matching and 2) camera parameters

$$\mathbf{x} = \alpha \mathbf{P}X$$

Same direction but differs by a scale factor

$$\mathbf{x} \times \mathbf{P}X = \mathbf{0}$$

Cross product of two vectors of same direction is zero
(this equality removes the scale factor)

Using the fact that the cross product should be zero

$$\mathbf{x} \times \mathbf{P}\mathbf{X} = \mathbf{0}$$

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \times \begin{bmatrix} p_1^\top \mathbf{X} \\ p_2^\top \mathbf{X} \\ p_3^\top \mathbf{X} \end{bmatrix} = \begin{bmatrix} yp_3^\top \mathbf{X} - p_2^\top \mathbf{X} \\ p_1^\top \mathbf{X} - xp_3^\top \mathbf{X} \\ xp_2^\top \mathbf{X} - yp_1^\top \mathbf{X} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Third line is a linear combination of the first and second lines.
(x times the first line plus y times the second line)

One 2D to 3D point correspondence give you 2 equations

$$\begin{bmatrix} y\mathbf{p}_3^\top \mathbf{X} - \mathbf{p}_2^\top \mathbf{X} \\ \mathbf{p}_1^\top \mathbf{X} - x\mathbf{p}_3^\top \mathbf{X} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} y\mathbf{p}_3^\top - \mathbf{p}_2^\top \\ \mathbf{p}_1^\top - x\mathbf{p}_3^\top \end{bmatrix} \mathbf{X} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\mathbf{A}_i \mathbf{X} = \mathbf{0}$$

Now we can make a system of linear equations
(two lines for each 2D point correspondence)

Concatenate the 2D points from both images

$$\begin{bmatrix} yp_3^\top - p_2^\top \\ p_1^\top - xp_3^\top \\ y'p_3'^\top - p_2'^\top \\ p_1'^\top - x'p_3'^\top \end{bmatrix} \mathbf{X} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\mathbf{A}\mathbf{X} = \mathbf{0}$$

How do we solve homogeneous linear system?

S V D !

Recall: Total least squares


(**Warning:** change of notation. \mathbf{x} is a vector of parameters!)

$$\begin{aligned} E_{\text{TLS}} &= \sum_i (\mathbf{a}_i \mathbf{x})^2 \\ &= \|\mathbf{A}\mathbf{x}\|^2 && \text{(matrix form)} \\ \|\mathbf{x}\|^2 &= 1 && \text{constraint} \end{aligned}$$

$$\begin{array}{ll} \text{minimize} & \|\mathbf{A}\mathbf{x}\|^2 \\ \text{subject to} & \|\mathbf{x}\|^2 = 1 \end{array} \quad \rightarrow \quad \begin{array}{ll} \text{minimize} & \frac{\|\mathbf{A}\mathbf{x}\|^2}{\|\mathbf{x}\|^2} \\ & \text{(Rayleigh quotient)} \end{array}$$


Solution is the eigenvector
corresponding to smallest eigenvalue of
 $\mathbf{A}^\top \mathbf{A}$

Depth from triangulated points



- Get camera coordinates of a point
- Take z-axis value -> depth!
- We need at least two views to get depth values (unless we rely on learning-based modules)

(u, v, d)
img
depth



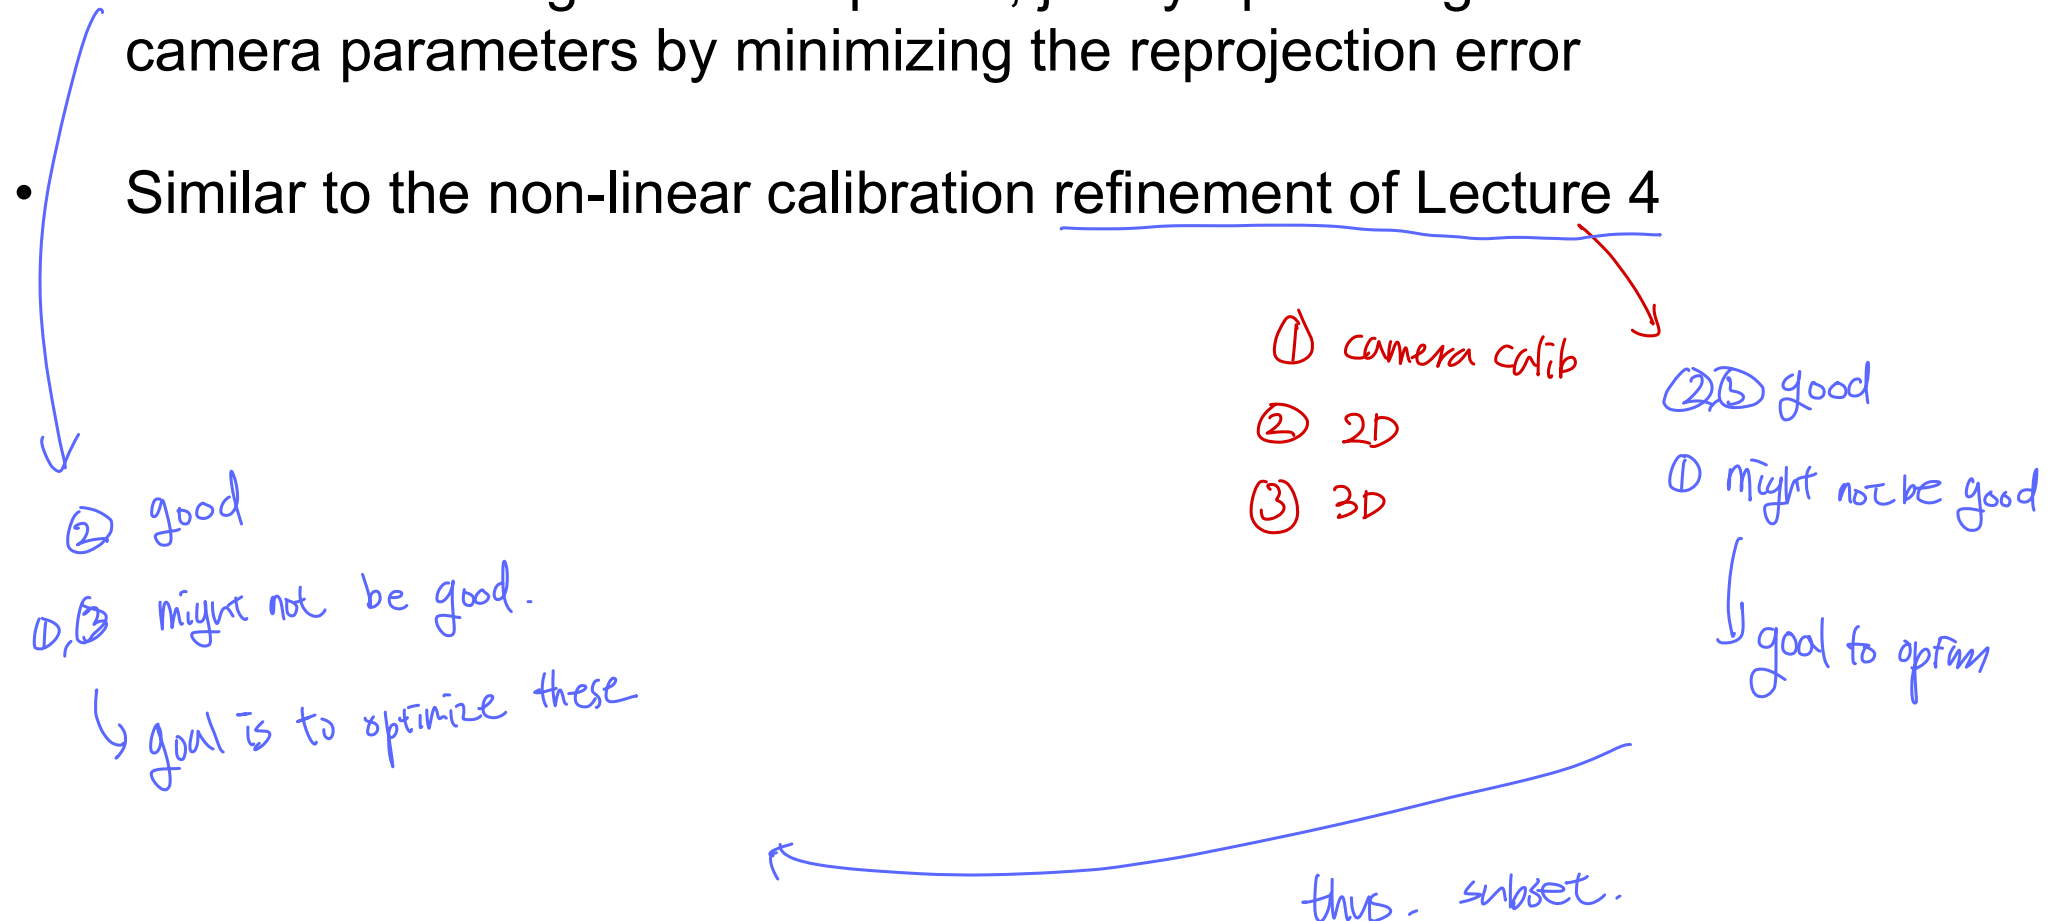
$X_w \rightarrow X_{C1}$
 $X_w \rightarrow X_{C2}$

Depth map
RGBD
H
W

2. Bundle Adjustment

Bundle Adjustment

- From initial triangulated 3D points, jointly optimizing 3D coordinates and camera parameters by minimizing the reprojection error
- Similar to the non-linear calibration refinement of Lecture 4




Reprojection Error

camera calibration

Slide from Lecture 4

- The reprojection error is the **Euclidean distance** (in pixels) between an **observed image point** and the **corresponding 3D point reprojected** onto the camera frame.
- The reprojection error gives us a **quantitative measure of the accuracy** of the calibration (**ideally it should be zero**).




Non-Linear Calibration Refinement

Slide from Lecture 4

- The calibration parameters K, R, T determined by the DLT can be refined by minimizing the following cost:

$$K, R, T, \text{ lens distortion} = \underset{K, R, T, \text{ lens}}{\operatorname{argmin}} \sum_{i=1}^n \|p^i - \pi(P_W^i, K, R, T)\|^2$$

- This time we also include the **lens distortion** (can be set to 0 for initialization)
- Can be minimized using **Levenberg–Marquardt** (more robust than Gauss-Newton to local minima)



● Control points
(observed points)

● Reprojected points
 $\pi(P_W^i, K, R, T)$


Non-Linear Calibration Refinement

Slide from Lecture 4

- The calibration parameters K, R, T determined by the DLT can be refined by minimizing the following cost:

$$K, R, T, \text{ lens distortion} = \underset{K, R, T, \text{ lens}}{\operatorname{argmin}} \sum_{i=1}^n \|p^i - \pi(P_W^i, K, R, T)\|^2$$

- This time we also include the **lens distortion** (can be set to 0 for initialization)
- Can be minimized using **Levenberg–Marquardt** (more robust than Gauss-Newton to local minima)



● Control points
(observed points)

● Reprojected points
 $\pi(P_W^i, K, R, T)$


Outlier Rejection

- There could be wrong correspondences across multi-view images -> wrong ones are called 'outliers'
- We should not use them for the bundle adjustment
- How can we reject outliers? -> Use RANSAC!

RANSAC

(**RAN**dom **SA**mples **C**onsensus) :

Fischler & Bolles in '81.




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example



Algorithm:


1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example

$$N_I = 6$$




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

keep repeat.

RANSAC



Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

How to choose parameters?

- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)
- Number of sampled points s
 - Minimum number needed to fit the model
- Distance threshold δ
 - Choose δ so that a good point with noise is likely (e.g., prob=0.95) within threshold
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$


$$N = \log(1-p) / \log(1-(1-e)^s)$$

no need to remember


just experimental result

RANSAC for 3D lifting

100 is ideal. but 5-10 is good for human.



Two hand touching range of motion



Algorithm:

1. Sample (randomly) the number of ^{view points.} points required for the triangulation
2. **Triangulate** the selected 2D points to the 3D space
3. Project the triangulated 3D points to all image space and check reprojection error. Reject viewpoints with huge error.

Repeat 1-3 until the best model is found with high confidence

Depth from triangulated points and bundle adjustment

- Get camera coordinates of a point with triangulation and bundle adjustment
- Take z-axis value -> depth!
- We need at least two views to get depth values (unless we rely on learning-based modules)

